

Handoff: Property-Management Partnership Playbook — First-Principles Strategy Review

TL;DR

- The playbook's core thesis — sell property managers a delegated "responsibility carried across organizational boundaries to a verified conclusion" service — is the **strongest of the three ideas Handoff has tested**, but it is a *services business with a software wedge*, not a venture-scale SaaS company. The honest venture paths are (a) an AI-implementation/operations-partner motion (fastest dollars) and (b), as the 12-month upside option, owning the P&L via an AI-native management firm (highest ceiling).
- The "adoption paradox" the founders observed is real but is **NOT caused by missing software**. It is caused by module-licensing/pricing gates, an implementation-and-configuration gap (Buildium's 2026 report found AI use jumped 20%→58% in a year while **only 8% of firms fully automated any workflow**), and cross-organizational responsibility boundaries that no vendor API can cross — all of which favor a permissioned human-in-the-loop operator, not another app.
- Move now and **concurrently**: ship an email/SMS/document-ingestion agent that produces a case queue + evidence-backed closeout packet in week one; sell paid diagnostic sprints and sign 1-2 paid Charlottesville student-housing pilots (Woodard, CBS, Management Services Corporation); and **price for outcomes** — \$15/case is far too low; target \$40-75/case or a \$1.25-3/door/month retainer. First \$10k in ~30 days is realistic. The prior rounds' "August 2027" and "manual phase first" defaults were wrong and should be discarded.

Key Findings

1. **Software is installed but incomplete, not absent.** The CBS Rentals RealPage case (ActiveBuilding portal + OneSite, yet manual reconstruction of repair/utility records) is representative, not anomalous. At industry scale, Buildium's 2026 *State of the Property Management Industry Report* (3,200+ professionals surveyed alongside NARPM) found AI adoption rose from 20% to 58% in a single year but **only 8% had fully automated any workflow**. The gap between "licensed" and "acted-on/verified" *is* the market.

2. **Pricing and module gates push work back onto humans and residents.** AppFolio's real floor is a 50-unit minimum plus an unpublished minimum spend; its AI (Realm-X Flows) and read/write API sit on the Plus tier (~\$3.00–3.20/unit, \$960 monthly minimum). EliseAI is quote-only, reported at \$3–6/unit/month with a ~\$25k annual minimum, viable only above several hundred units. Small and mid operators buy the system of record and skip the orchestration/AI modules — so the capability "exists" but is never bought or configured.
 3. **The residual work is a boundary problem, not a feature problem.** No API can grant access to a resident's private utility account, decide a contested obligation, authorize spend outside a management agreement, or force a city/utility/vendor to respond. That is exactly the "connective tissue" the playbook names — and it is genuinely unowned in many shops.
 4. **PMs already outsource this shape of work, which de-risks the thesis.** APM Help (per-unit PM accounting/bookkeeping including move-outs and owner payments; competitor OJO bills \$34–45/hr), Latchel (~\$25/unit/mo maintenance coordination, or resident-paid at \$15–18/mo), Property Meld (\$1.60–2.00/unit/mo), (Aitoolsbakery) Second Nature (resident benefits, \$20–80/mo resident-paid), (Second Nature) and offshore PM VAs/BPOs (\$6–25/hr; ShoreAgents full-time \$1,200–2,500/mo) prove a real buyer exists for delegated closeout and coordination.
 5. **The venture-shape question is unresolved and must be faced head-on.** AI-enabled service roll-ups are the hottest 2025–26 category (General Catalyst's Creation Fund allocated ~\$1.5B; its Long Lake HOA-management platform raised ~\$670M and reached \$100M EBITDA in under two years; Thrive Holdings; 8VC). But the skeptics are loud: in *Fortune* (June 27, 2025) Air Street's Nathan Benaich and PolyAI's Nikola Mrkšić called the "services trade like software" arbitrage "a mirage... confusing operational improvement with business model transformation," noting BPOs trade at 5–23x EV/EBITDA versus software's 22–92x.
 6. **Staffing pain is a genuine tailwind.** The 2025 RCLCO survey (reported by NAA, April 2026) put onsite property-level turnover at **29.2%** (33.4% at 450+-employee orgs; voluntary turnover 23.4%, up from 21.7% in 2024); 74% of PM professionals cite HR/staffing/recruitment among their top-three challenges. **Capacity relief — not headcount cuts — is the sellable benefit.**
 7. **Licensing is navigable but not trivial.** Virginia's §54.1-2103 exemptions are narrow (a licensed broker's employee may show units, give factual info, take applications, accept broker-approved documents, and accept deposits/rent — not negotiate leases for compensation for another). NY RPL §440/§440-a defines "collecting rent for another for a fee" as brokerage requiring a license. Handoff must operate as the manager's *delegate* under the manager's license and (optionally) brand — not as an independent unlicensed manager.
-

Details

PART A — First-principles diagnosis of the "adoption paradox"

Evidence — module licensing and pricing gates.

- **AppFolio:** Core ~\$1.40–1.49/unit/mo with a **50-unit minimum and \$298 monthly minimum**; Plus ~\$3.00–3.20/unit with a **\$960 minimum**; Max ~\$5/unit with a \$1,500–7,500 minimum. (Appfoliopricing) At exactly 50 units the effective Core rate is ~\$5.60–5.96/unit (roughly 4× headline). Realm-X AI Flows and read/write API are gated to Plus/Max. (managecasa.com, costbench.com, usecarly.com, itqlick.com; verified 2026.)
- **Buildium:** Essential ~\$62/mo, Growth ~\$192/mo, Premium ~\$400/mo, each covering ~150 units before door scaling, plus heavy transaction fees (\$2.35 incoming EFT on Essential). (Checkthat) The Lumina AI suite is Premium-tier. (BoardStack) (buildium.com/pricing, checkthat.ai, renpro.com; 2026.)
- **EliseAI:** quote-only, reported **\$3–6/unit/mo with a ~\$25k annual minimum** — a figure that circulates only in competitor comparisons (aitoolsbakery, realty-ai, meltflexai) and should be treated as *reported, not confirmed*. It pencils only above several hundred units. EliseAI itself reports (Businesswire, June 2026) **\$200M ARR** with "100% year-over-year growth for five consecutive years," powering "one in eight apartments in the United States"; its \$250M Series E (Aug 2025) was led by a16z at a **\$2.2B valuation**, having crossed \$100M ARR.

Inference: The capability chain "Available → suitable → licensed → connected → configured → acted-on → verified" breaks for small/mid operators at *licensed* and *configured*. The saturated market is saturated with *record-keeping*, not *execution*.

Evidence — the implementation/configuration gap. Buildium 2026: 20%→58% AI adoption, only **8% fully automated any process**. AppFolio's 2026 benchmark (1,617 US residential PM pros, surveyed Sept–Nov 2025): ~44–45% use AI, 30% cite utility setup as a top move-in problem, and a large share plan tech consolidation. *Forbes* (July 8, 2026) framed the paradox precisely: "if the economics are this compelling, why hasn't the industry simply completed the transition?" — because most platforms "sit on top of existing operations, making staff faster rather than redundant." **This is the single most important datapoint for Handoff: the gap between AI use and AI completion is the wedge.**

Evidence — incentive misalignment. The resident suffers but the operator buys; shifting coordination onto residents is cheap for managers (CBS asking the resident for water-usage history; the NYC broker email routing the resident across super/portal/utility/insurance/access channels). Owner-operators (CBS) internalize resident pain differently from third-party

managers, who owners judge on fees and vacancy, not closeout quality. Buildium's May 2026 survey of 300 small owners (95% own 1-20 units): 84% want approval over large repairs; communication/service quality drive switching.

Evidence — fragmentation of responsibility. The utility leak-credit case crosses resident → management → utility → city, each with its own authority and no single owner of the outcome. This is *organizational*, not technical, fragmentation — the part no API fixes.

Evidence — NYC structure. The FARE Act (Local Law 119 of 2024, effective **June 11, 2025**) shifted landlord-hired broker fees off tenants; (Cole Schotz) landlords now absorb roughly one month's rent in leasing cost per unit and have fresh incentive to cut turnover friction and leasing spend. NYC ops stay email/text/phone-driven because responsibility is split across management companies, 32BJ supers, brokers, owners, and co-op/condo boards, with RPL §440 gating rent collection "for another." Family-owned 500–5,000-unit firms fit the profile: real workflows, thin tech configuration.

Evidence — Charlottesville owner-operators. Woodard Properties (~700 UVA students, (Woodard Properties) ~\$7M revenue, (RocketReach) operating since 1981) (Woodard Properties) and CBS (RealPage stack) show that even a RealPage-equipped owner-operator leaves closeout gaps — because the software is a *system of record*, not an *accountable operator of cross-party cases*.

Evidence — labor. RCLCO 2025 (via NAA, April 2026): 29.2% onsite turnover (33.4% at large orgs), voluntary turnover 23.4% up from 21.7%. AI has layered onto staff, not displaced them (the 8% full-automation figure).

Diagnosis — which frictions are real, addressable, or structural:

- *Real + addressable by Handoff:* incomplete configuration/execution; unowned cross-party closeout; capacity relief for thin/turnover-prone teams; documentation and verification.
- *Real but structural (not Handoff-fixable):* private account access, contested-obligation decisions, third-party non-response, licensing boundaries. Handoff can *coordinate and escalate* these but cannot *own the decision*.
- *Partly a pricing artifact:* small operators skip modules — a reason to *sell an operated outcome*, not another license.

PART B — Is "responsibility to a verified conclusion" a fundable business?

Challenge 1: Is "nobody is responsible" true, or is responsibility deliberately excluded because it's unprofitable? The evidence cuts both ways. APM Help proves administrative-to-financial closeout is a paid market; Latchel proves "AI + humans to completion" already sells. But the utility-leak-credit archetype is low-frequency and high-variance — the playbook itself

concedes "a handful of difficult cases can dominate the economics." **Verdict:** responsibility is unowned in the middle band — too complex to fully automate, too infrequent to staff a specialist, too cross-org for the PMS — but the *profitable* slice is the repeatable closeout (deposits, roommate changes, renewals reconciliation), not the pathological one-offs.

Challenge 2: Precedent. AI-enabled service roll-ups define the 2025–26 venture category — General Catalyst's Creation Fund (~\$1.5B), including Long Lake (HOA management, ~\$670M raised, \$100M EBITDA in under two years); Thrive Holdings; 8VC; with "property management" explicitly named as a target vertical. Direct comparables: **Doorstead** (tech-enabled full-service PM, founded 2019 by Ryan Walianny and Jennifer Bronzo, ~\$37M total raised through a \$21.5M Series B, 5% monthly fee, upfront rent guarantee) and Belong. AI-native PM back-office rounds in 2025–26 include **PARES AI** (YC S2025, ~\$500K), **Lette AI** (\$1.4M pre-seed), and **Rely** (\$4.5M seed, May 2026, 2048 Ventures) for multifamily diligence. The counter-precedent is the *Fortune* critique above: operational improvement ≠ business-model transformation, and BPO multiples (5–23x) sit far below software (22–92x).

Challenge 3: Do PMs buy "managed workflow delegation"? Yes — in narrow, well-defined lanes: accounting (APM Help; OJO \$34–45/hr), maintenance coordination (Latchel ~\$25/unit; Property Meld \$1.60–2.00/unit), (Aitoolsbakery) call centers, and VAs (\$6–25/hr). They do **not** buy vague "responsibility." Handoff must sell a *named, bounded case type with an SLA and an evidence packet*, priced like APM/Latchel — not a philosophy.

Challenge 4: Legal boundary.

- *Virginia (§54.1-2103):* exemptions are narrow — an unlicensed party may show units, give factual info, take applications, accept broker-approved documents, and accept deposits/rent, (Virginia Code) **only as the employee of the licensed broker managing the property**. Negotiating lease terms for compensation for another requires a license; unlicensed activity is a Class 1 misdemeanor (Management Pro) with fines up to \$1,000/transaction (§54.1-2105.2). (TenantCloud)
- *New York (RPL §440/§440-a):* "collecting rent for another for a fee" and negotiating rentals is brokerage requiring a license; (New York State Senate) §440-a exemptions are minimal.
- *Implication:* Handoff must operate as the manager's *delegate* under the manager's license and (optionally) brand, with a written scope, delegated-access-only, and no independent handling of contested funds or lease negotiation. Workable — it is how VAs/BPOs already operate — but it caps how "customer-facing" Handoff can be and forbids the unlicensed-independent-manager path.

Verdict on B: Stronger than the prior two theses (unit-centric transition layer; student-housing transition case) because it (a) has a named payer with budget precedent, (b) sits on the structurally-unautomatable side of the 8% line, and (c) is defensible via trusted delivery and permissioned integrations rather than a copyable feature. But it is a **services business with a software wedge**, and its venture ceiling depends on *reuse across partners*, not on the closeout insight itself.

PART C — The two lanes, ranked, plus better ones

Lane 1 — Approved household changes / roommate swaps in shared student rentals.

- *Frequency:* High in student portfolios; CBS's published renewal process distinguishes roommate changes from turns, requires multi-party consent, carries deposits forward. Replacement/reassignment routinely triggers admin, re-key, and application fees (LawInsider clause samples; anecdotal ~\$300 name-swap fees).
- *Incumbent overlap:* Rent Manager, Rent Manager Yardi Student, Entrata Student, and AppFolio (student housing at \$1.40/unit) already support multi-occupant leases and roommate management — the *software* exists; the *cross-party closeout execution* (chasing signatures, reconciling effective dates, deposit hand-off between residents, account updates) is unowned.
- *Economic stake:* modest per case but recurring and seasonal; managers may already charge fees Handoff could operate and grow.
- *Delegability:* high — bounded, policy-driven, non-decisional.
- **Rank: strong wedge, low ACV.**

Lane 2 — Repair-to-billing closeout (incl. utility-leak-credit).

- *Frequency:* continuous but overlaps Property Meld/HappyCo/Latchel/APM Help heavily.
- *Economic stake:* high variance; the leak-credit archetype is infrequent.
- *Delegability:* medium (touches contested charges).
- **Rank: weaker; crowded and variance-heavy.**

Better/additional lanes (first-principles):

- **Deposit dispositions as compliance-critical closeout — the best net-new lane.** Virginia (Va. Code §55.1-1226) mandates an itemized return within **45 days** (a single 15-day extension only if a third-party contractor is needed and the tenant is

notified in writing within the 45-day window); willful non-compliance exposes the landlord to actual damages and attorney fees, and missing the deadline can forfeit the right to *any* deductions. New York (GOL §7-108) mandates **14 days**, with automatic forfeiture of any right to retain the deposit if missed, the landlord bearing the burden of proving reasonableness, and **punitive damages of up to twice the deposit** for willful violation (per *Colon v. Martin*, 35 N.Y.3d 75). This is high-frequency (every move-out), deadline-driven, penalty-backed, evidence-heavy, and cross-party (maintenance photos + accounting + resident comms) — exactly Handoff's shape, with a *quantifiable cost of failure*.
Rank: #1 on economic stake × frequency × delegability.

- **Utility transfer/verification at move-in/out** (AppFolio's 30%-cite move-in pain) — high frequency, bounded.
- **Insurance/renters-policy compliance chasing** — recurring, rules-based.
- **NYC violation/HPD-compliance chasing and super coordination** — high stakes, cross-org, deadline-driven.
- **Vendor invoice reconciliation / owner-reporting escalation** — overlaps APM Help.

Lane ranking (frequency × economic stake × delegability × low incumbent overlap × speed to first dollar):

1. Deposit disposition/compliance closeout (VA 45-day / NY 14-day-with-double-damages)
2. Roommate-change closeout (student)
3. Utility transfer/verification
4. NYC HPD/violation + super coordination
5. Repair-to-billing closeout
6. Insurance compliance chasing

PART D — Partnership architecture

Model	Speed to revenue	Ceiling	Dependency risk	Founder fit
D1. Delegated operator FOR PM companies (playbook)	Fast (weeks)	Medium — services margins, reuse-limited	Medium (per-partner scope)	High — quant/AI + ops
D2. AI implementation/config partner for incumbents (fill the 8% gap)	Fastest (days-weeks)	Medium — consulting-like but recurring	High (platform dependency)	Very high — software/AI build
D3. Broker/brokerage trigger partner (NYC)	Medium	Low-medium	High (FARE-disrupted channel)	Medium
D4. White-label / co-sell with a point vendor	Medium	Medium	High	Medium
D5. Own the P&L — AI-native PM firm / rollup	Slow (quarters)	Highest — software multiple only if reuse proven	Low (own the customer)	High if paired with capital

- **D2 is the fastest cash and the most under-appreciated:** the 8%-full-automation figure is a last-mile implementation market. LeadSimple already runs a paid consultant/certified-partner ecosystem (Black Sheep Global, PM PathBuilders, Real-Time Consulting, etc.); AppFolio Stack, Yardi, and Entrata (which announced 100+ embedded agentic AI agents in March 2026) run marketplaces. An "AI implementation partner that actually *finishes* configuration and then *operates* the workflow" can bill setup + retainer immediately, with real discovery baked in. Risk: platform dependency and commoditization by the vendors' own CS teams.
- **D1 is the playbook's model** and the best *product-discovery* vehicle — best combined with D2 (get paid to implement, stay to operate).
- **D5 is the highest ceiling** and the honest answer to "incumbents can trivially copy the feature": if the moat is trusted execution, owning the P&L (an AI-native student-housing manager in Charlottesville, or a small NYC book) captures full

margin and compounds. Doorstead (~\$37M raised, 5% fee) shows the model attracts venture capital; the risk is it is operations-heavy and slow.

Recommendation on architecture: Run **D2+D1 concurrently now** (implementation-to-operation partner with 1-2 PMs), and hold **D5 as the 12-month option** if reuse and margins prove out. The two-person AI-native team's edge (fast shipping, quant rigor, evaluation discipline) fits D2/D1 build-and-operate far better than a capital-intensive rollout on day one.

PART E — Economics, rebuilt

Is \$1,800 for 120 cases (= \$15/case) sane? No — it is too low. Benchmarks: Latchel ~\$25/unit/mo for coordinated resolution; (Sfailabs) APM per-unit bookkeeping; VAs \$6-25/hr; OJO \$34-45/hr. (APM Help) A compliance-grade deposit disposition or a multi-party roommate closeout is worth far more than \$15 to a manager facing a 14-day New York forfeiture with double-damages exposure. **Target \$40-75/case** for compliance closeout, or a **per-door retainer** (\$1.25-3/door/mo) plus overage, or **cohort pricing** for seasonal student turns.

Marginal cost in 2026 with an AI-native stack: the playbook's 10-min/case human-review assumption is *pessimistic* for repeatable cases once templates and evaluations exist — LLM extraction/comparison/drafting can take routine cases to 2-5 min of human verification. It is *optimistic* for the pathological long tail (contested charges, non-responsive third parties), which can consume 25+ min and turn contribution negative. Blend accordingly and price the tail out (exclusions, overage, escalation back to the manager).

Three scenarios (illustrative):

Scenario	Model	Price	Gross margin	1,500-unit ACV	5,000-unit ACV
Service-heavy	Humans do 60%+	\$50/case, ~1 case/unit/yr	30-40%	~\$75k	~\$250k
Hybrid (target)	AI 80%, human verify	\$45/case or \$2/door/mo	55-65%	~\$36k-70k	~\$120k-250k
Software-led	AI 95%, exception-only human	\$1.25/door/mo + overage	70-80%	~\$22k	~\$75k

Path to ARR (hybrid): \$1M ~ 15-25 mid-size partners or 5-8 large ones; \$5M ~ 40-80 partners or a few enterprise books + multi-lane expansion; \$20M requires either enterprise multifamily penetration (competing with EliseAI's distribution) or the

D5 owned-P&L flip. **Honest read:** at software-led margins with proven reuse this is a fundable seed/Series A "services-as-software" story; at service-heavy margins it is a good cash business VCs will discount to BPO multiples (the Fortune/Air Street critique). Investors *have* funded this exact shape in 2025-26 (Creation Fund's Long Lake; Doorstead; PARES/Lette/Rely) — but the ones that flip to software multiples are those that demonstrate the margin move on their *own first customers* before asking investors to underwrite scale.

PART F — Speed and sequencing (30/60/90, concurrent)

Week 1 (build + sell in parallel):

- **Ship:** an inbound email/SMS/document-ingestion agent that parses a resident/manager thread plus attachments (leases, work orders, utility bills) against **exported** PMS data (CSV from AppFolio/Buildium/RealPage — no integration needed), produces a **case queue** (next action / accountable party / evidence / time-waiting), and generates an **evidence-backed closeout packet** (PDF). Days of work for an AI-native pair, not months.
- **Sell:** book discovery+pilot calls with **CBS Rentals, Woodard Properties, Management Services Corporation, Wade Apartments, KRS Property Management, PMI Commonwealth** and other Charlottesville student operators; in NYC, target family-owned 500-5,000-unit managers and small third-party firms. Use the founders' lived CBS/NYC experience as the wedge.
- **Charge from day one:** a paid **diagnostic sprint** (\$2,500-5,000) reconstructing 20-30 consecutive deposit or roommate-change cases and quantifying the manager's current cost — simultaneously discovery and revenue.

Day 30: 1-2 signed paid pilots on a single lane (deposit closeout or roommate closeout), scoped in writing (case types, approval thresholds, channels, hours, source of truth, error handling, retention, handback), operating under the manager's brand/license as delegate. **First \$10k booked.**

Day 60: live cases running with scorecard tracking (cases completed correctly, partner effort saved, reopened cases, missing approvals, duplicate/unauthorized actions, elapsed vs active time); add a second partner to force the reuse test; open a parallel D2 implementation engagement with one PM to fund the build.

Day 90 (success state): ≥ 2 paying partners; one lane delivered at $>55\%$ gross margin; a reusable case definition + onboarding + evaluation set; and a decision on whether to expand lanes, add partners, or begin the D5 owned-P&L pilot. **First-\$100k path:** 3-5 partners on retainer + diagnostic sprints, or one enterprise book on a per-door retainer, over ~90-120 days.

Risks the aggressive path accepts: (1) operating before a signed scope is legally sloppy — get the delegation agreement first; (2) export-based (non-integrated) data means manual refresh and staleness risk — fine for a pilot, not for scale; (3) concurrent build+sell+operate on a two-person team risks quality lapses on the pathological tail — cap case types and exclusions hard; (4) compliance lanes carry real penalty exposure — Handoff coordinates and documents, but the manager must retain the sign-off.

PART G — Kill/pivot conditions and the counter-thesis

Steelman that the partnership thesis is ALSO wrong:

- PMs may refuse to delegate resident-facing, funds-adjacent work to an unlicensed third party (VA/BPO precedent suggests they will — but only for back-office, non-decisional tasks).
- The residual "boundary" work may be low-value *by design*: managers exclude it because it doesn't pay, structurally capping willingness-to-pay (the playbook concedes WTP and margins are unproven).
- EliseAI + Latchel + APM Help already sell "AI + humans to completion" across leasing, maintenance, and accounting — Handoff's insight may be a feature they add, not a company.
- The founders' pain is partly a *consumer* pain (resident experience) with no reliable B2B payer.
- The correct move may be to **own the operator** (D5), or to leave PropTech for a market where the founders' quant/AI edge compounds faster (the PM-software TAM is modest — ~\$6.5B in 2026 per Mordor Intelligence).

Kill/pivot triggers:

- **Kill the delegated-service thesis** if, after 5–8 discovery reconstructions, managers cannot name a case type they will pay >2× fully-costed delivery for, or refuse written delegation.
- **Pivot to D2 (implementation)** if PMs pay for configuration/operation but won't delegate ongoing case ownership.
- **Pivot to D5 (own the P&L)** if execution quality is the proven moat and a Charlottesville student book is acquirable/operable.
- **Leave** if neither WTP nor reuse appears by ~90–120 days across ≥3 partners.

Final recommendation (decisive):

1. **Commit to the delegated-operator + implementation-partner motion (D1+D2) immediately, one lane, deposit-disposition closeout first** — compliance-backed WTP; VA 45-day / NY 14-day-with-double-damages makes the cost of failure quantifiable.
 2. **Ship the ingestion → queue → closeout-packet agent in week one** against exported PMS data; do not wait for integrations.
 3. **Sell paid diagnostic sprints from day one** to named Charlottesville operators; convert to scoped pilots within 30 days; book first \$10k in ~30 days, first \$100k in ~90-120 days.
 4. **Price for outcomes, not messages:** \$40-75/case or \$1.25-3/door/mo + overage; exclude/escalate the pathological tail.
 5. **Instrument everything** (the playbook's scorecard) and force the reuse test with a second partner by day 60.
 6. **Hold D5 (own the P&L) as the 12-month option;** pursue it only if execution-as-moat and margins are proven — the honest answer to "incumbents can copy the feature."
 7. **Reject the prior rounds' conservative defaults** (token manual phase, August-2027 timelines, "do the service manually first"): the fast version runs discovery, build, and paid pilots concurrently, and the only non-negotiable gate is the signed delegation scope before touching live cases.
-

Recommendations

Execute the seven-step plan in Part G: a concurrent D1+D2 motion, deposit-disposition closeout as the beachhead lane, week-one shipping against exported data, paid diagnostic sprints converting to scoped pilots, outcome-based pricing, aggressive instrumentation and a forced reuse test by day 60, and the owned-P&L firm (D5) held as the 12-month upside option. Benchmarks that would change the plan: if ≥ 3 partners won't sign written delegation or name a $>2\times$ -cost case type within ~90-120 days → kill/pivot; if PMs pay to configure but not to delegate → pivot to D2; if execution quality proves to be the moat and a student book is operable → accelerate D5.

Caveats

- EliseAI's ~\$25k minimum and \$3-6/unit pricing are *reported by competitors*, not confirmed by EliseAI; treat as directional. (EliseAI's \$200M ARR, ~1-in-8-apartments, and \$2.2B Series E valuation are company/press-confirmed.)
- The playbook's unit economics are illustrative; real WTP and margins are unproven and must be validated in the first 30-60 days.

- Roommate-change and deposit-case *frequencies* are inferred from lease structures and statutes, not from a measured Charlottesville/NYC dataset — the diagnostic sprint exists to measure them.
- The legal analysis here is not legal advice; confirm the delegation structure with Virginia and New York counsel before operating live cases. Note the "up to double the deposit" Virginia penalty is less consistently sourced than the confirmed actual-damages-plus-attorney-fees remedy; the New York double-damages/forfeiture provisions are confirmed in the statute text.
- The services-vs-software venture multiple is genuinely contested (Fortune/Air Street critique vs. Creation Fund thesis); the outcome depends on demonstrated reuse and margin, not on the closeout insight alone.